

Xiaocong Yang

PhD Candidate CS @ UIUC | Founder, [AI Interpretability @ Illinois](#)

[Google Scholar](#) | [LinkedIn](#) | xy51@illinois.edu | Urbana-Champaign

Education

Ph.D. in Computer Science University of Illinois Urbana-Champaign	2025.8 - 2028.5
- Founder and current leader of <i>AI Interpretability @ Illinois</i> research team.	
M.S. in Computer Science University of Illinois Urbana-Champaign	2023.8 - 2025.5
B.S. in Economics Tsinghua University	2018.9 - 2023.6

Publication

6. **Xiaocong Yang**, Jiacheng Lin, Ziqi Wang, Chengxiang Zhai. *Learning by Analogy: Enhancing Few-Shot Prompting for Math Word Problem Solving with Computational Graph-Based Retrieval*, preprint [[link](#)]
5. Ziyi Chen, **Xiaocong Yang**, Jiacheng Lin, Chenkai Sun, Jie Huang, Kevin Chen-Chuan Chang. *Cascade Speculative Drafting for Even Faster LLM Inference*, **NeurIPS 2024** [[link](#)]
4. **Xiaocong Yang**, James Huang, Wenxuan Zhou, Muhao Chen. *Parameter-efficient Tuning with Special Token Adaptation*, **EACL 2023** [[link](#)]
3. Xingcheng Yao, Yanan Zheng, **Xiaocong Yang**, Zhilin Yang. *NLP from Scratch Without Large-Scale Pretraining: A Simple and Efficient Framework*, **ICML 2022** [[link](#)]
2. Leixian Shen, Enya Shen, Yuyu Luo, **Xiaocong Yang**, Xuming Hu, Xiongshuai Zhang, Zhiwei Tai, Jianmin Wang. *Towards natural language interfaces for data visualization: A survey*, **TVCG 2022** [[link](#)]
1. Hao Zhou, Pei Ke, Zheng Zhang, Yuxian Gu, Yinhe Zheng, Chujie Zheng, Yida Wang, Chen Henry Wu, Hao Sun, **Xiaocong Yang**, Bosi Wen, Xiaoyan Zhu, Minlie Huang, Jie Tang. *Eva: An open-domain chinese dialogue system with large-scale generative pre-training*, preprint [[link](#)]

Research Experience

AI Scientist Intern Intuit Inc.	2025.5-2025.8
<i>Supervisor: Ruocheng Guo Staff Research Scientist, Intuit AI Research</i>	Mountain View
• Raised RestBench tool-call success by +5 percentage points by designing tool description improvement pipeline for agentic LLM flows; scoped and led end-to-end study at Intuit.	
Student Researcher TIMAN Group, University of Illinois Urbana-Champaign	2024.3 - present
<i>Advisor: Chengxiang Zhai Donald Biggar Willett Professor, UIUC</i>	Urbana-Champaign
• Proposed and led research on <i>Learning by Analogy</i> , a RAG system in which the retriever is trained to retrieve examples with similar computational graphs to enhance few-shot prompting by providing the LLM with the correct reasoning path. Our approach outperforms the baseline by +6.7 percentage points .	
Student Researcher FORWARD Data Lab, University of Illinois Urbana-Champaign	2023.8 - 2024.1
<i>Advisor: Kevin Chenchuan Chang Professor, UIUC</i>	Urbana-Champaign
• Proposed <i>Cascade Speculative Drafting</i> , where we vertically and horizontally cascade multiple draft models of different size. Our approach is mathematically proved to be more efficient than standard speculative decoding and achieves up to 72% additional speedup in the experiments.	
• Paper is published at NeurIPS 2024.	
Visiting Researcher LUKA Group, University of Southern California	2022.4 - 2022.9
<i>Advisor: Muhao Chen Assistant Professor, Information Science Institute, USC</i>	Los Angeles
• Proposed and led research on <i>PASTA</i> , a model finetuning method that is more efficient than all of existing baselines which tunes only 0.029% of total parameters to match the performance of full finetuning.	
• Illustrated special tokens' role as information disseminators in LLMs.	
• Paper is published at EACL 2023.	

Student Researcher | IIIS, Tsinghua University

2021.7 - 2023.1

Advisor: [Zhilin Yang](#) | Assistant Professor, Institute for Interdisciplinary Information Sciences, Tsinghua University | Founder & CEO, Moonshot AI Beijing

- Contributed to *Task-driven Language Modeling*, proposing a task-specific pretraining framework that reduced by **100x** the total training cost and size of corpus and democratizing the development of large LLMs.
- Analyzed the functionality of “vertical attention heads” in LLMs.
- Paper is published at ICML 2022.

Student Researcher | Conversational AI (CoAI) Group

2021.1 - 2021.6

Advisor: [Huang Minlie](#) | Associate Professor | Department of Computer Science and Technology, Tsinghua University Beijing

- Contributed to the development of the open-domain dialogue system *EVA* and the corresponding paper.
- Designed the interface and client-server communication framework for the deployment of the dialog system, tested rule-based and model-based methods to filter out toxic content in the generated text.

Teaching Experience

-
- Instructor** | CS 591 BAI, UIUC 2025.8 - present
 - Teaching Assistant** | CS441 Applied Machine Learning, UIUC 2024.1 - 2025.5
 - Teaching Assistant** | CS107 Data Science Discovery, UIUC 2023.8 - 2023.12

Community Experience

-
- Student Member** | oSTEM @ Illinois 2023.11 - present